Hypothesis quality and cross-model hypothesis composition agreement
All models averaged (3 datasets)

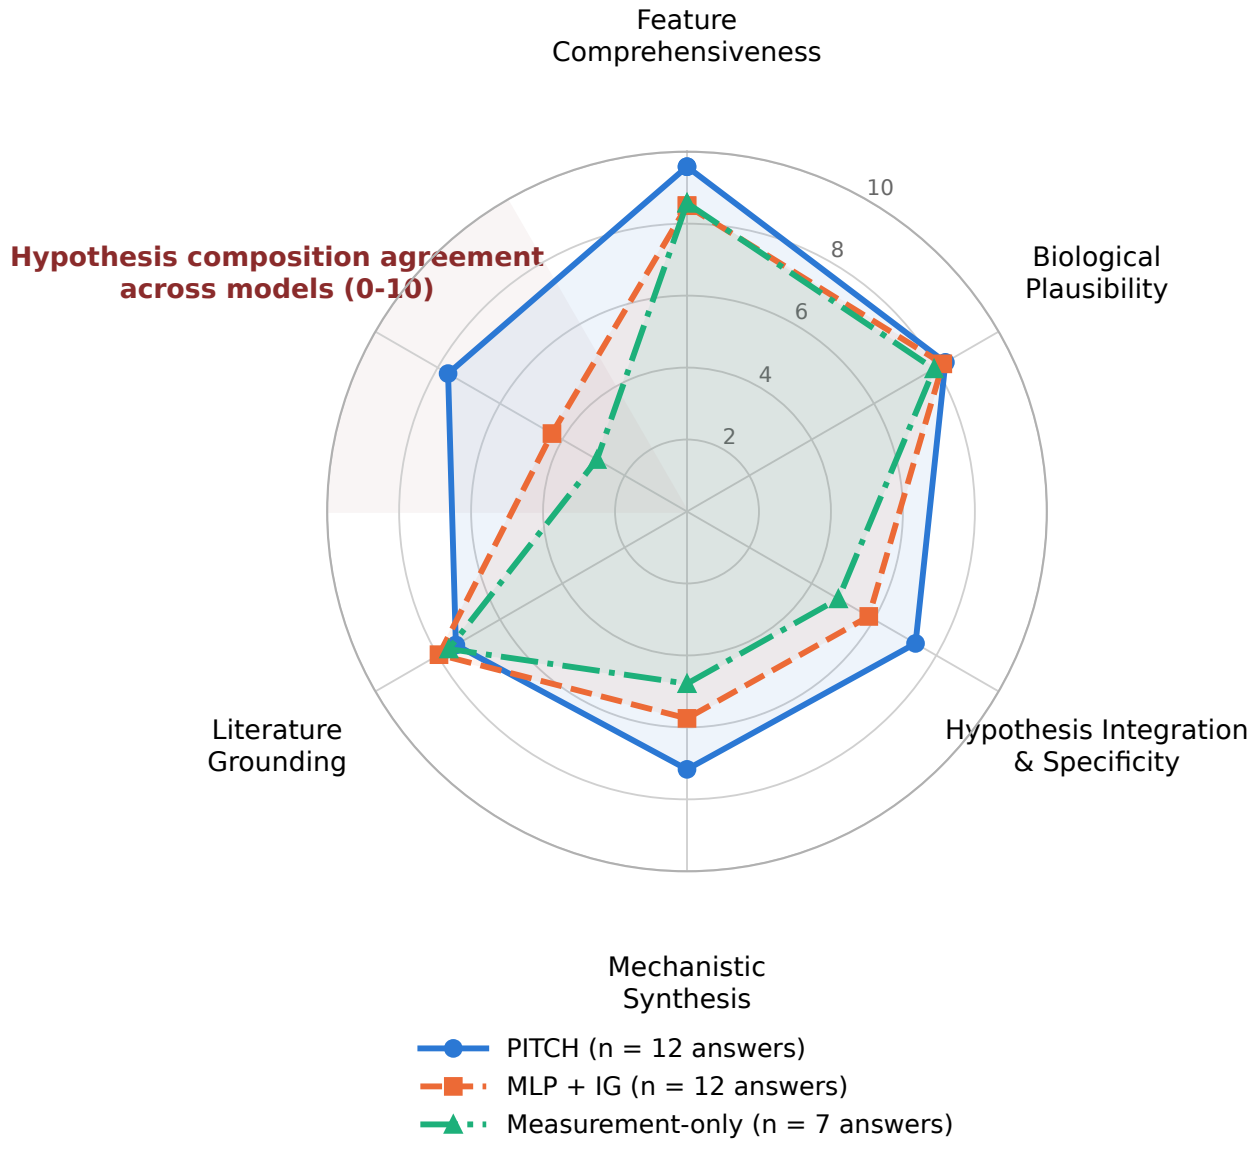

Measurement-only: one committed answer per available (model, prompt) cell, 7 of 12 cells (GLM 3, Qwen 2, DeepSeek 1, Gemini 1); PITCH and MLP + IG: 12 answers each.
Axis 6: agreement across the 4 models for PITCH and MLP + IG; for measurement-only, among the models with a successfully generated answer (3, 2 and 2 answers per prompt).