Hypothesis quality and cross-model hypothesis composition agreement
All models averaged (3 datasets)

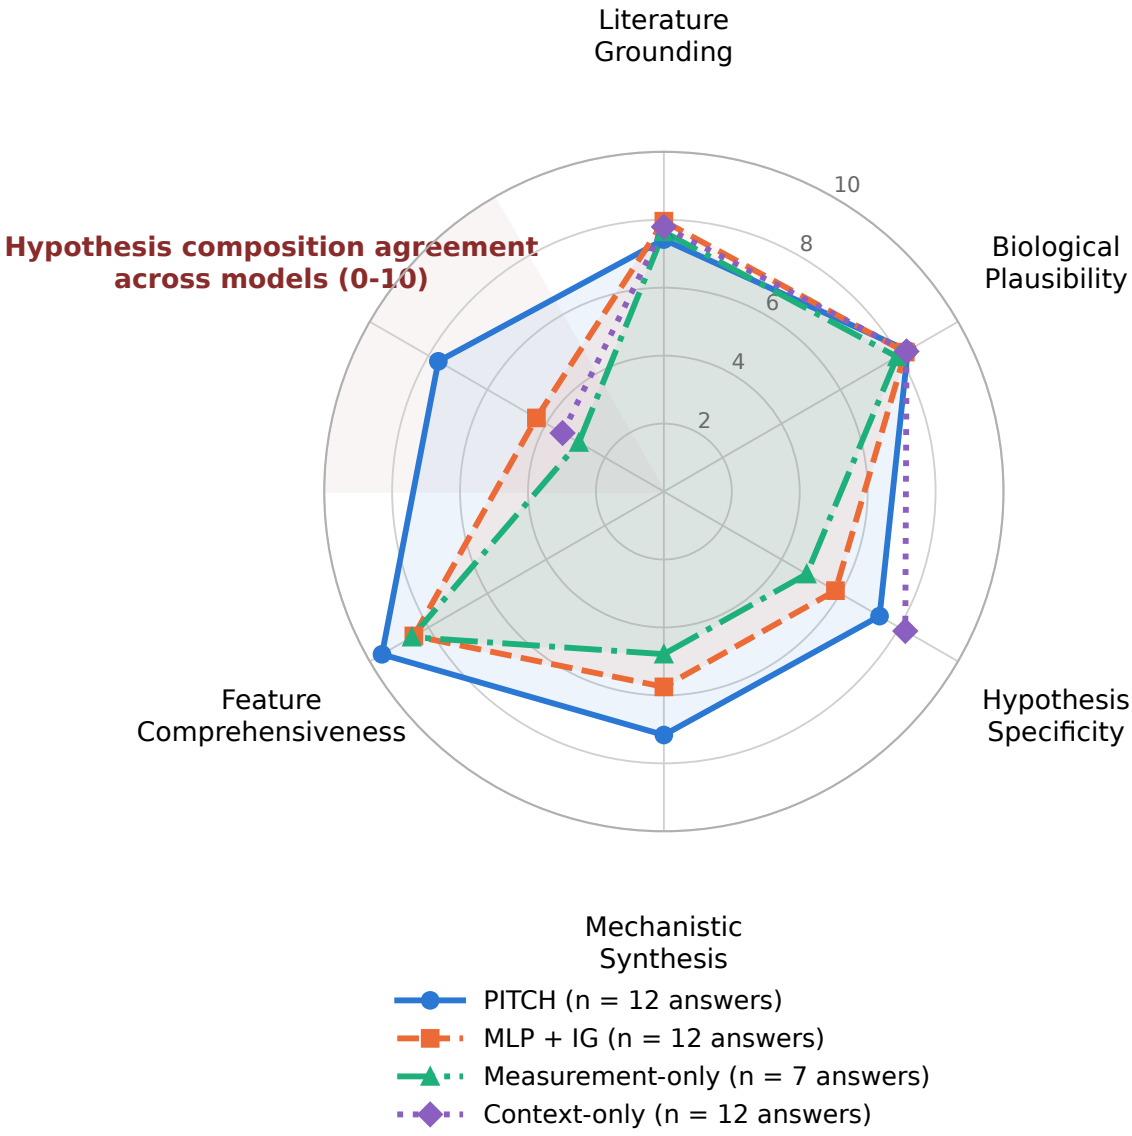

n = successfully generated answers (one random committed answer per model and prompt):
PITCH 12, MLP + IG 12, Measurement-only 7, Context-only 12.
Axis 6: PITCH and MLP + IG across all 4 models; Measurement-only and Context-only among the models with a successfully generated answer for the prompt (answers per prompt, CITE-seq / CosMx / Gouwens: Measurement-only 2, 3, 2; Context-only 4, 4, 4).
Context-only has no value on Feature Comprehensiveness and Mechanistic Synthesis (these criteria are defined on the supplied focus / contextual features and do not apply when no data are given).