Rubric scores per model, grouped by similarity category
22 models, 6 cases

	Tree reading		Feature analysis	Biological validity	Hypothesis construction		Supplementary
	Zone identification	Rule extraction (logic rigor)	Feature comprehensiveness	Biological plausibility	Hypothesis integration & specificity	Mechanistic synthesis	Literature grounding (bonus, half weight)
Mistral-Small-3.1-24B	9.2	9.3	8.8	6.0	3.7	3.0	4.0
Llama-3.3-70B-Instruct-Turbo	9.2	9.3	8.3	7.7	4.2	3.7	5.7
phi4-14B	7.5	8.3	8.7	6.2	3.8	3.7	4.8
Qwen3.5-9B-think-temp1	10.0	9.3	10.0	4.2	7.7	4.3	3.3
ministral3-14B	9.2	9.3	9.8	5.2	6.8	5.7	3.3
Nemotron3-Nano-30B-A3B	9.2	7.0	9.8	3.8	7.8	4.7	2.7
Qwen3-Next-80B-A3B-Instruct	10.0	9.3	10.0	4.2	8.5	5.2	3.3
gemma-4-12B-it	9.2	9.3	9.8	6.3	6.7	5.3	4.3
gemma-4-31B-it	9.2	8.8	10.0	7.7	7.0	6.5	5.8
Qwen3.8-27B-think	9.2	9.3	10.0	5.3	8.2	5.8	3.7
DeepSeek-V4.1-Flash	9.2	9.3	10.0	8.8	7.8	7.5	8.2
gpt-oss-120b	10.0	9.7	9.7	5.5	8.0	6.3	3.5
DeepSeek-V4-Flash-0731	9.2	9.3	10.0	8.3	8.0	8.0	7.8
DeepSeek-V4-Pro-0813	10.0	10.0	10.0	8.5	8.0	7.8	7.7
MiniMax-M3	10.0	10.0	10.0	7.0	8.3	7.5	5.5
Qwen3.8-2.4T-A95B	10.0	10.0	10.0	8.8	8.5	7.8	7.8
GLM-5.3-Flash	9.2	9.0	10.0	8.7	7.8	7.5	7.3
Kimi-K3	9.2	9.3	10.0	8.2	8.7	8.5	8.0
GLM-5.3	9.2	9.3	10.0	8.3	9.0	8.8	7.5
gemini-3.1-pro-preview	9.2	9.3	9.8	8.2	7.8	6.7	6.2
Qwen3.6-35B-A3B-think	10.0	10.0	10.0	5.3	8.5	6.3	4.0
Inkling	10.0	10.0	10.0	6.8	8.3	8.0	5.8

Category (color)
- Category 1: Rule-faithful descriptive annotators
- Category 2: Integrative, weakly grounded hypothesis generators
- Category 3: Mechanistic, literature-grounded integrators

Mean rubric score (0-10)
10
8
6
4
2
0

Each cell = model's mean over 6 cases. Row order matches the similarity dendrogram heatmap.